

In [50]: `import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf`

In [51]: `train = pd.read_csv(r'/Users/kvno1ahmednagar/Desktop/DATASETS/fashionmnist-2/fashion-mnist_train.csv')
test = pd.read_csv(r'/Users/kvno1ahmednagar/Desktop/DATASETS/fashionmnist-2/fashion-mnist_test.csv')`

In [52]: `from sklearn.model_selection import train_test_split`

In [53]: `train.head()`

Out[53]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777	pixel778	pixel779
0	2	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
1	9	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
2	6	0	0	0	0	0	0	0	5	0	...	0	0	0	30	4
3	0	0	0	0	1	2	0	0	0	0	...	3	0	0	0	0
4	3	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows × 785 columns

In [54]: `train.shape`

Out[54]: (60000, 785)

In [55]: `test.shape`

Out[55]: (10000, 785)

In [56]: `train_data = np.array(train, dtype = 'float32')
test_data = np.array(test, dtype = 'float32')`

In [57]: `X_train = train_data[:,1:]
Y_train = train_data[:,0]`

In [13]: `X_train.head()`

AttributeError Traceback (most recent call last)
<ipython-input-13-73020798007e> in <module>
----> 1 X_train.head()

AttributeError: 'numpy.ndarray' object has no attribute 'head'

In [58]: `X_train.shape`

Out[58]: (60000, 784)

In [59]: `Y_train.shape`

Out[59]: (60000,)

In [60]: `X_test = test_data[:,1:]
Y_test = test_data[:,0]`

In [61]: `X_test.shape`

Out[61]: (10000, 784)

In [62]: `Y_test.shape`

Out[62]: (10000,)

In [63]: `X_train, X_valid, Y_train, Y_valid = train_test_split(X_train, Y_train, test_size = 0.2)`

In [64]: `X_train.shape`

Out[64]: (48000, 784)

In [65]: `Y_train.shape`

Out[65]: (48000,)

In [66]: `X_valid.shape`

Out[66]: (12000, 784)

In [67]: `Y_valid.shape`

Out[67]: (12000,)

In [68]: `rows = 28
cols = 28
im_shape = (rows, cols, 1)

X_train = X_train.reshape(X_train.shape[0], *im_shape)
X_test = X_test.reshape(X_test.shape[0], *im_shape)
X_valid = X_valid.reshape(X_valid.shape[0], *im_shape)`

In [69]: `X_train.shape`

Out[69]: (48000, 28, 28, 1)

In [70]: `from keras.utils import np_utils`

In [76]: `Y_train = np_utils.to_categorical(Y_train, num_classes = 10)
Y_valid = np_utils.to_categorical(Y_valid, num_classes = 10)
Y_test = np_utils.to_categorical(Y_test, num_classes = 10)`

In [71]: `import keras
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, BatchNormalization, Flatten, Dense, Dropout
from keras.optimizers import Adam`

In [72]: `from keras.regularizers import l2`

In [73]: `model = Sequential()
model.add(Conv2D(filters = 64, kernel_size = (3,3), strides = (1,1), activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPooling2D(pool_size = (3,3), strides = (1,1)))
model.add(Dropout(0.25))

model.add(Conv2D(filters = 32, kernel_size = (3,3), strides = (1,1), activation = 'relu'))
model.add(MaxPooling2D(pool_size = (2,2), strides = (1,1)))
model.add(Dropout(0.15))

model.add(Conv2D(filters = 128, kernel_size = (5,5), strides = (1,1), activation = 'relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size = (3,3), strides = (1,1)))
model.add(Dropout(0.20))

model.add(Flatten())
model.add(Dense(128, activation = 'relu'))
model.add(Dense(10, activation = 'softmax'))`

In [74]: `model.compile(loss = 'categorical_crossentropy', optimizer = Adam(lr = 0.001), metrics = ['accuracy'])`

In [77]: `model.fit(X_train, Y_train, batch_size = 128, epochs = 15, verbose = 1, validation_data = (X_valid, Y_valid))`

WARNING:tensorflow:From /opt/anaconda3/lib/python3.7/site-packages/keras/backend/tensorflow_backend.py:422: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.

Train on 48000 samples, validate on 12000 samples

Epoch 1/15
48000/48000 [=====] - 1091s 23ms/step - loss: 1.8568 - accuracy: 0.7469 - val_loss: 1.5247 - val_accuracy: 0.7204
Epoch 2/15
48000/48000 [=====] - 923s 19ms/step - loss: 0.4347 - accuracy: 0.8491 - val_loss: 0.5780 - val_accuracy: 0.8310
Epoch 3/15
48000/48000 [=====] - 1407s 29ms/step - loss: 0.3520 - accuracy: 0.8740 - val_loss: 0.5425 - val_accuracy: 0.8340
Epoch 4/15
48000/48000 [=====] - 536s 11ms/step - loss: 0.3233 - accuracy: 0.8816 - val_loss: 0.3778 - val_accuracy: 0.8747
Epoch 5/15
48000/48000 [=====] - 561s 12ms/step - loss: 0.3002 - accuracy: 0.8916 - val_loss: 0.3413 - val_accuracy: 0.8820
Epoch 6/15
48000/48000 [=====] - 2268s 47ms/step - loss: 0.2767 - accuracy: 0.8994 - val_loss: 0.3778 - val_accuracy: 0.8746
Epoch 7/15
48000/48000 [=====] - 637s 13ms/step - loss: 0.2601 - accuracy: 0.9034 - val_loss: 0.3369 - val_accuracy: 0.8877
Epoch 8/15
48000/48000 [=====] - 1037s 22ms/step - loss: 0.2529 - accuracy: 0.9086 - val_loss: 0.3056 - val_accuracy: 0.8931
Epoch 9/15
48000/48000 [=====] - 1553s 32ms/step - loss: 0.2353 - accuracy: 0.9130 - val_loss: 0.3707 - val_accuracy: 0.8860
Epoch 10/15
48000/48000 [=====] - 554s 12ms/step - loss: 0.2223 - accuracy: 0.9183 - val_loss: 0.2741 - val_accuracy: 0.9053
Epoch 11/15
48000/48000 [=====] - 620s 13ms/step - loss: 0.2200 - accuracy: 0.9194 - val_loss: 0.3559 - val_accuracy: 0.8822
Epoch 12/15
48000/48000 [=====] - 589s 12ms/step - loss: 0.2072 - accuracy: 0.9231 - val_loss: 0.2797 - val_accuracy: 0.9045
Epoch 13/15
48000/48000 [=====] - 548s 11ms/step - loss: 0.2010 - accuracy: 0.9263 - val_loss: 0.3091 - val_accuracy: 0.9003
Epoch 14/15
48000/48000 [=====] - 548s 11ms/step - loss: 0.1951 - accuracy: 0.9286 - val_loss: 0.2799 - val_accuracy: 0.9109
Epoch 15/15
48000/48000 [=====] - 544s 11ms/step - loss: 0.1887 - accuracy: 0.9295 - val_loss: 0.3239 - val_accuracy: 0.8982

Out[77]: <keras.callbacks.callbacks.History at 0x64780bed0>

In [78]: `score = model.evaluate(X_test, Y_test, verbose = 0)
print('Test Loss: {:.4f}'.format(score[0]))
print('Test Acc: {:.4f}'.format(score[1]))`